

DI WU

+1 (310)498-8193 ✧ xiaowu200031@gmail.com ✧ <https://xiaowu0162.github.io/>

EDUCATION

University of California, Los Angeles

Los Angeles, CA, USA

Ph.D. Candidate in Computer Science (Advisor: Prof. Kai-Wei Chang)

09/2022 - present

- Overall GPA: **4.000/4.000**
- Research Interests: long-horizon agents, retrieval-augmented language models, text generation evaluation.

University of California, Los Angeles

Los Angeles, CA, USA

B.S. in Computer Science, Summa Cum Laude

08/2018 - 06/2022

- Overall GPA: **3.973/4.000**

RESEARCH EXPERIENCE

UCLA NLP Lab

Los Angeles, CA, USA

Student Researcher, Supervisor: Prof. Kai-Wei Chang

02/2021 - present

- Research on neural keyphrase generation methods and their evaluation.
- Research on building robust retrieval-augmented generation and long-term memory methods.
- Collaborate with Taboola on a number of real-world systems including concept extraction, continual learning, knowledge agents, and automatic content moderation.
- Collaborate with Prof. Tian Yuan's team on building robust coding agents for Amazon Nova Challenge.

The Ozcan Research Group, UCLA

Los Angeles, CA, USA

Student Research Assistant, Supervisor: Prof. Aydogan Ozcan

09/2019 - 06/2022

- Researched on virtual staining of skin tissues with an emphasis on Basal Cell Carcinoma. Designed methods to improve both single-image quality and temporal coherence of the predictions.
- Explored dataset engineering and data augmentation methods to mitigate class imbalance.

PUBLICATIONS

1. **Di Wu**, Zixiang Ji, Asmi Kawatkar, Bryan Kwan, Jia-Chen Gu, Nanyun Peng, and Kai-Wei Chang, 2026. LongMemEval-V2: Evaluating Long-Term Agent Memory Toward Experienced Colleagues. *preprint, ICML 2026 SCALE Workshop (Oral)*.
2. **Di Wu**, Devendra Singh Sachan, Wen-tau Yih, and Mingda Chen, 2026. Procedural Knowledge at Scale Improves Reasoning. *COLM 2026*.
3. **Di Wu**^{*}, Yixin Wan^{*}, and Kai-Wei Chang, 2025. VisRet: Visualization Improves Knowledge-Intensive Text-to-Image Retrieval. *ACL 2026 (Oral)*.
4. **Di Wu**^{*}, Jia-Chen Gu^{*}, Kai-Wei Chang, and Nanyun Peng, 2025. Self-Routing RAG: Binding Selective Retrieval with Knowledge Verbalization. *preprint*.
5. Jia-Chen Gu, Junyi Zhang, **Di Wu**, Yuankai Li, Kai-Wei Chang, and Nanyun Peng, 2025. BRIEF-Pro: Universal Context Compression with Short-to-Long Synthesis for Fast and Accurate Multi-Hop Reasoning. *Findings of the ACL: ACL 2026*.
6. **Di Wu**, Siyue Liu, Zixiang Ji, Ya-Liang Chang, Zhe-Yu Liu, Andrew Pleffer, and Kai-Wei Chang, 2025. Open-Domain Safety Policy Construction. *Findings of the ACL: EACL 2026*.
7. Yiming Wang, Da Yin, Yuedong Cui, Ruichen Zheng, Zhiqian Li, Zongyu Lin, **Di Wu**, Xueqing Wu, Chenchen Ye, Yu Zhou, and Kai-Wei Chang, 2025. LLMs as Scalable, General-Purpose Simulators for Evolving Digital Agent Training. *NeurIPS 2025 LAW Workshop*.
8. Yuankai Li, Jia-Chen Gu, **Di Wu**, Kai-Wei Chang, and Nanyun Peng, 2024. BRIEF: Bridging Retrieval and Inference via Multi-hop Reasoning and Compression. *Findings of the ACL: NAACL 2025*.

9. **Di Wu**, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu, 2024. LongMemEval: Benchmarking Chat Assistants on Long-Term Interactive Memory. *ICLR 2025*.
10. **Di Wu**, Jia-Chen Gu, Fan Yin, Nanyun Peng, and Kai-Wei Chang, 2024. Synchronous Faithfulness Monitoring for Trustworthy Retrieval-Augmented Generation. *EMNLP 2024*.
11. Yixin Wan, **Di Wu**, Haoran Wang, and Kai-Wei Chang, 2024. The Factuality Tax of Diversity-Intervened Text-to-Image Generation: Benchmark and Fact-Augmented Intervention. *EMNLP 2024*.
12. **Di Wu**, Xiaoxian Shen, and Kai-Wei Chang, 2024. MetaKP: On-Demand Keyphrase Generation. *Findings of the ACL: EMNLP 2024*.
13. **Di Wu**, Da Yin, and Kai-Wei Chang, 2024. KPEval: Towards Fine-grained Semantic-based Keyphrase Evaluation. *Findings of the ACL: ACL 2024*.
14. **Di Wu**, Wasi Uddin Ahmad, Dejiao Zhang, Murali Krishna Ramanathan, and Xiaofei Ma, 2023. Repoformer: Selective Retrieval for Repository-level Code Completion. *ICML 2024 (Oral)*.
15. **Di Wu**, Wasi Uddin Ahmad, and Kai-Wei Chang, 2023. On Leveraging Encoder-only Pre-trained Language Models for Effective Keyphrase Generation. *LREC-COLING 2024*.
16. Yuzhu Li, Nir Pillar, Jingxi Li, Tairan Liu, **Di Wu**, Songyu Sun, Guangdong Ma, Kevin de Haan, Luzhe Huang, Yijie Zhang, Sepehr Hamidi, Anatoly Urisman, Tal Keidar Haran, William Dean Wallace, Jonathan E. Zuckerman, and Aydogan Ozcan, 2024. Virtual histological staining of unlabeled autopsy tissue. *Nature Communications*, 15(1), p.1684.
17. **Di Wu**, Wasi Uddin Ahmad, and Kai-Wei Chang, 2023. Rethinking Model Selection and Decoding for Keyphrase Generation with Pre-trained Sequence-to-Sequence Models. *EMNLP 2023*.
18. Po-Nien Kung, Fan Yin, **Di Wu**, Kai-Wei Chang, and Nanyun Peng, 2023. Active Instruction Tuning: Improving Cross-Task Generalization by Training on Prompt Sensitive Tasks. *EMNLP 2023*.
19. **Di Wu**, Wasi Uddin Ahmad, and Kai-Wei Chang, 2023. Pre-trained Language Models for Keyphrase Generation: A Thorough Empirical Study. *Preprint*.
20. **Di Wu**, Wasi Uddin Ahmad, Sunipa Dev, and Kai-Wei Chang, 2023. Representation Learning for Resource-Constrained Keyphrase Generation. *Findings of the ACL: EMNLP 2022*.
21. Jingxi Li, Jason Garfinkel, Xiaoran Zhang, **Di Wu**, Yijie Zhang, Kevin de Haan, Hongda Wang, Tairan Liu, Bijie Bai, Yair Rivenson, Gennady Rubinstein, Philip O. Scumpia, and Aydogan Ozcan, 2023. Biopsy-free in vivo virtual histology of skin using deep learning. *Light: Science & Applications*, 10(1), 1-22.

INTERNSHIP EXPERIENCE

Google Cloud AI Research

Student Researcher; Mentors: Rujun Han, Yanfei Chen, Chen-Yu Lee

Sunnyvale, CA, USA

06/2026 - 09/2026

- Researching on continually evolving multi-agent systems.

FAIR at Meta Superintelligence Labs

Research Scientist Intern; Mentors: Mingda Chen, Devendra Sachan, Scott Yih

New York, NY, USA

06/2025 - 10/2025

- Researched on memory representations for improving frontier reasoning tasks.
- Built the first large-scale procedural memory datastore that significantly improves various reasoning models.

Tencent AI Lab

Research Intern; Mentors: Hongwei Wang, Wenhao Yu

Bellevue, WA, USA

06/2024 - 09/2024

- Researched on long-term memory for chat assistants and built a large-scale high-quality evaluation benchmark.
- Designed indexing strategies for improving long-term memory retrieval performance.

AWS AI Labs

Applied Scientist Intern; Mentors: Wasi Ahmad, Dejiao Zhang

New York, NY, USA

06/2023 - 09/2023

- Researched on improving retrieval-augmented code language models for repository-level code completion.
- Formulated the task of selective retrieval-augmented infilling. Designed approaches from the perspective of in-repository code retrievers and the code generator models.

- The designed model achieves no performance loss with only 10% of retrieval budget.

Microsoft Research Asia

Research Intern; Mentor: Ning Shang

Beijing, China
04/2021 - 09/2021

- Researched on model compression and AutoML. Contributed to the open source project [NNI](#).

NewsBreak

Natural Language Processing Intern

Beijing, China
07/2020 - 10/2020

- Worked on a hierarchical multi-label classification problem with 268 categories. Improved the f1-score by 49% with multiple statistical and deep learning methods.
- Improved the performance of model pre-training, fine-tuning, and online serving pipelines.

TEACHING

Teaching Assistant

- UCLA CS 33, Introduction to Computer Organization, Fall 2023, with Prof. Glenn Reinman.
- UCLA CS 33, Introduction to Computer Organization, Spring 2024, with Prof. Glenn Reinman.
- UCLA CS 33, Introduction to Computer Organization, Fall 2024, with Prof. Tony Nowatzki.

SERVICES AND AWARDS

- Reviewer: ACL 2023, EMNLP 2023, AAAI, ICLR, ICML, NeurIPS, JAIR, ACL Rolling Review.
- Outstanding Reviewer, EMNLP 2024.
- Amazon Trainium Fellowship, Fall 2025.
- UCLA HSSEAS Dean's List, 8 times, 2018-2022.